

Parallel programming in Chapel

ALEX RAZOUMOV
alex.razoumov@westgrid.ca



Why another language? <https://chapel-lang.org>

- High-level parallel programming language
 - ▶ “Python for parallel programming”
 - ▶ much easier to use and learn than MPI; few lines of Chapel typically replace tens of lines of MPI code
 - ▶ abstractions for data distribution/parallelism, task parallelism
 - ▶ optimization for data-driven placement of subcomputations
 - ▶ granular (“multi-resolution”) design: can bring closer to machine level if needed
 - ▶ everything you can do in MPI (and OpenMP!), you should be able to do in Chapel
- Shared- and distributed-memory parallelism built in
- Focus on performance
 - ▶ compiled language; simple Chapel codes perform as well as optimized C/C++/Fortran code
 - ▶ reportedly, very complex Chapel codes run at ~70% performance of a similar well-tuned MPI code (not bad, but room to improve)
- Perfect language for learning parallel programming for beginners
- Open-source: can compile on any Unix-like platform
 - ▶ precompiled for MacOS (single-locale via Homebrew)
 - ▶ Docker image <http://dockr.ly/2vJbi06> (simulates a multi-locale environment)

- Cedar (OmniPath) / Graham (InfiniBand) / Béluga (InfiniBand)

▶ <https://docs.computecanada.ca/wiki/Chapel>

```
$ source /home/razoumov/startSingleLocale.sh
```

```
$ chpl --version
```

```
$ source /home/razoumov/startMultiLocale.sh
```

```
$ chpl --version
```

- Fairly small community at the moment:
too few people know/use Chapel \iff relatively few libraries
- You can use functions/libraries written in other languages, e.g. in C
 - (1) direct calls will always be serial
 - (2) high-level Chapel parallel libraries can use C/F90/etc libraries underneath

Useful links

- Slides from <https://chapel-lang.org>
 - ▶ *Data parallelism*
 - ▶ *Task parallelism*
 - ▶ *Locality / Affinity Features*
 - ▶ *Domain Maps / Distributions*
- Watch *Chapel: Productive, Multiresolution Parallel Programming* talk by Brad Chamberlain
- *Getting started guide for Python programmers*
- <https://learnxinyminutes.com/docs/chapel>
- Concise *Chapel tutorial* by David Bunde
- Documentation and examples for various Chapel modules in `$CHPL_HOME/modules/`, e.g., `standard/` or `dists/`
- <https://stackoverflow.com/questions/tagged/chapel>

Our workshop

PART 1: BASIC LANGUAGE FEATURES

- running single-locale Chapel codes on Cedar
 - ▶ interactive jobs vs. batch jobs
- quicky on running Chapel on your laptop
- problem description: heat transfer equation
- variables
- ranges and arrays
- conditionals
- `for` loops
- config variables
- timing code execution

PART 2: TASK PARALLELISM

- parallel concepts
 - ▶ concurrency vs. true parallelism
 - ▶ concurrency vs. task locality
- fire-and-forget tasks
 - ▶ `begin` statement
 - ▶ `cobegin` statement
 - ▶ `coforall` loops
 - ▶ `forall` loops
- task synchronization
 - ▶ `sync` statement
 - ▶ sync variables
 - ▶ atomic variables
- task-parallelizing the heat transfer solver (if we have time)

PART 3: DOMAIN PARALLELISM

- running multi-locale Chapel codes on Cedar
- simple multi-locale codes
- domains and single-locale data parallelism
- distributed domains
- heat transfer solver on distributed domains
- periodic boundary conditions
- writing to files

Numerical problem: 2D heat transfer equation

- Imagine a metallic plate initially at 25 degrees
- Simple 2D heat (diffusion) equation

$$\frac{\partial T(x, y, t)}{\partial t} = \frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2}$$

- Discretize the solution $T(x, y, t) \approx T_{i,j}^{(n)}$ with $i = 1, \dots, \text{rows}$ and $j = 1, \dots, \text{cols}$
 - ▶ upper left corner is $(1, 1)$, lower right corner is $(\text{rows}, \text{cols})$
- Initial condition: $T_{i,j}^{(0)} = 25$
- Boundary condition: upper side $T_{0,1..\text{cols}}^{(n)} \equiv 0$, left side $T_{1..\text{rows},0}^{(n)} \equiv 0$,
bottom side $T_{\text{rows}+1,1..\text{cols}}^{(n)} = 80 \cdot j/\text{cols}$, right side $T_{1..\text{rows},\text{cols}+1}^{(n)} = 80 \cdot i/\text{rows}$
(linearly increasing from 0 to 80 degrees)
- Discretize the equation with forward Euler time stepping

$$\frac{T_{i,j}^{(n+1)} - T_{i,j}^{(n)}}{\Delta t} = \frac{T_{i+1,j}^{(n)} - 2T_{i,j}^{(n)} + T_{i-1,j}^{(n)}}{(\Delta x)^2} + \frac{T_{i,j+1}^{(n)} - 2T_{i,j}^{(n)} + T_{i,j-1}^{(n)}}{(\Delta y)^2}$$

Numerical problem: 2D heat transfer equation (cont.)

- For simplicity assume $\Delta x = \Delta y = 1$
- Use $\Delta t = 1/4$
- The finite difference equation becomes

$$T_{i,j}^{(n+1)} = \frac{1}{4} \left[T_{i+1,j}^{(n)} + T_{i-1,j}^{(n)} + T_{i,j+1}^{(n)} + T_{i,j-1}^{(n)} \right]$$

- The objective is to find $T_{i,j}$ after a certain number of iterations, or when the system is in steady state
- Can increase the number of points in the grid to illustrate the advantage of parallelism


Serial exercise: using *procedures* and control flow

Look up Chapel procedures


Write a Chapel code to find the root of the equation

$x^5 + 8x^3 - 2x^2 + 5x - 1.2 = 0$ using the bisection method in the interval $[-1,1]$

- Calculate the function at the ends and the midpoint of the interval
- Depending on the signs of the three computed values, let the midpoint be either the new left or the new right end
- Repeat until your error is below $\Delta x = 10^{-8}$




Parallelism vs. TASK LOCALITY



Task- vs. data-parallel

	single locale shared memory parallelism	multiple locales distributed memory parallelism + likely shared memory parallelism
task parallel	<pre>config var numtasks = 2; coforall taskid in 1..numtasks do writeln("this is task ", taskid);</pre>	<pre>forall loc in Locales do on loc do writeln("this locale is named ", here.name);</pre>
data parallel	<pre>var A, B, C: [1..1000] real; forall (a,b,c) in zip(A,B,C) do c = a + b;</pre>	<pre>use BlockDist; const mesh = {1..100,1..100} dmapped Block(boundingBox={1..100,1..100}); var T: [mesh] real; forall (i,j) in T.domain do T[i,j] = i + j;</pre>

Array decomposition



Race condition

```
lock.add(1)  
lock.waitFor(2)
```



"task 1 is done"

```
lock.sub(1)
```

```
lock.waitFor(0)
```



•
•
•

```
lock.add(1)
```

```
lock.waitFor(2)
```



Note: `lock.waitFor()` is not a collective operation

Data-parallel exercise: compute π with `forall` loop



Write a parallel Chapel code to compute π by calculating the integral numerically through summation

$$\pi = \int_0^1 \frac{4 dx}{1 + x^2}$$

Parallelism cheatsheet

- `for` is a serial loop; `a..#n` means `n` iterations, `a..b` means `b-a+1` iterations
- `forall` loop is executed cooperatively by all local cores in parallel, or by remote locales that own the corresponding indices/elements (subdividing their local iterations among their local cores); number of threads scales to the number of available cores
- `coforall` loop creates a new task per each iteration (cycling through locales or tasks inside a locale)
- `begin { ... }` spins statements inside off into a new task
- `sync { ... }` pauses until the children have synced back up
- `cobegin { line1 line2 line3 }` runs each line in a new task; can be grouped with `{ }`
- Built-in variables and arrays
 - ▶ `numLocales` is the number of locales
 - ▶ `Locales` stores an array of compute nodes on which the program is executing
 - ▶ `locale.id` is the ID of the current locale
 - ▶ `locale.maxTaskPar` is the runtime maximum number of tasks on the current local
 - ▶ `locale.numCores` is the locale's number of compute cores
 - ▶ `locale.name` is a locale's name
 - ▶ `here` evaluates to the locale on which the current task is running
- Distributions
 - ▶ `BlockDist` partitions indices into blocks according to a boundingBox domain and maps each block onto a separate locale
 - ▶ `CyclicDist` maps indices to locales in a round-robin pattern starting at a given index
 - ▶ `BlockCycDist`, `DimensionalDist2D`, `PrivateDist`, `ReplicatedDist`, `StencilDist`, `BlockCycDim`, `BlockDim`, `ReplicatedDim`

Distributed domains



Unstructured data

In addition to rectangular domains and arrays, Chapel supports less structured (and more dynamic) data

- Sparse domains and arrays: can be mapped to locales
- Associative domains and arrays: can be mapped to locales as of v1.19
 - associated domains are similar to Python's sets (unordered, unique indices)
 - associated domains with arrays on top are similar to Python's dictionaries
 - `HashedDist` supports custom index mapping to locales
- Opaque domains and arrays: distribution across locales currently not implemented

WestGrid Chapel webinars

<https://westgrid.github.io/trainingMaterials/programming>

- Three-part “INTRO TO PARALLEL PROGRAMMING IN CHAPEL”
- “WORKING WITH DISTRIBUTED UNSTRUCTURED DATA IN CHAPEL”
- “WORKING WITH DATA FILES AND EXTERNAL C LIBRARIES IN CHAPEL”